

SQLBITS

Deploy Databricks Components using Terraform



 @annawykes

 <https://www.linkedin.com/in/anna-maria-wykes-31939454>

 AnnaWykes

 @falekmiah

 falekmiah.com

 FalekMiah01



advancinganalytics.co.uk

Anna Wykes

Senior Consultant

- 16+ Years Microsoft Data Analytics
- Intensive Software & Data Engineering Experience
- Microsoft MVP, Data Consultant & Public Speaker



Falek Miah

Principal Consultant

- 14+ Years Microsoft Data Analytics
- Intensive Data Engineering Experience
- Data, Cloud & DevOps Enthusiast
- Microsoft Azure, Databricks (Spark), Terraform (HashiCorp) certified





Session Scope

Session Scope

 Manually provisioning Databricks workspace and its components (clusters, libraries, etc) in the cloud can be:

- time consuming
- hard to repeat
- becomes inconsistent between environments



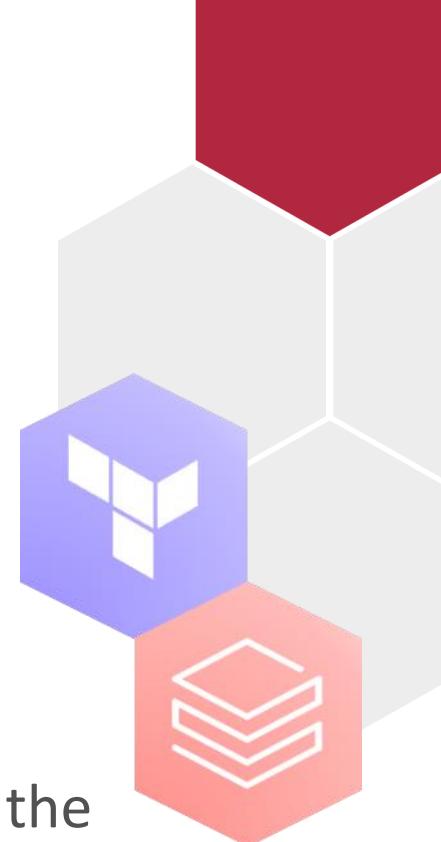
Using Infrastructure as Code (IaC) tool like Terraform can minimize this.



Databricks Labs introduce a collection of Terraform Providers that gives you the ability to deploy nearly all resources onto your cloud platform using Terraform.



Reduces the need to use PowerShell, Databricks CLI or REST APIs.

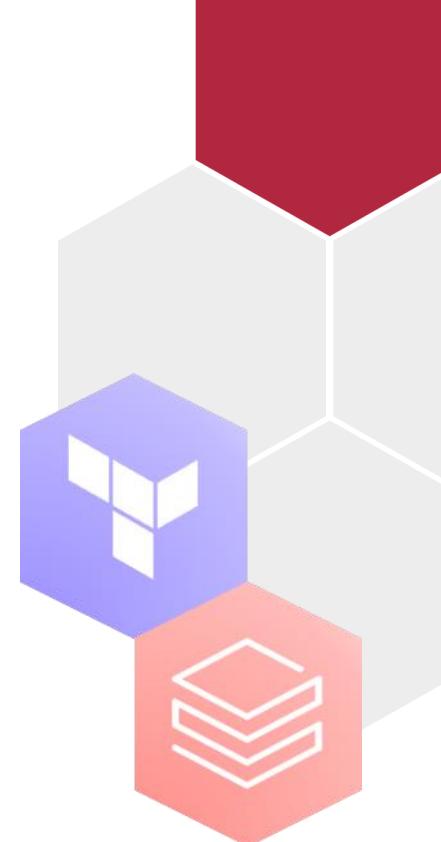




Infrastructure as Code (IaC)

Infrastructure as Code (IaC)

- **Automate** the creation, update or destroy of cloud infrastructure
- Allows you to **version**, **share**, and **reuse** cloud infrastructure
- **Configure** script (blueprint) of cloud infrastructure



Popular Infrastructure as Code (IaC) Tools



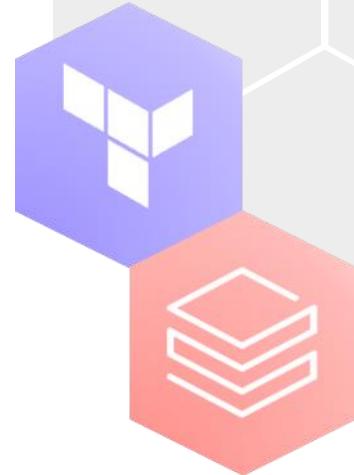
ARM (Azure Resource Manager)
Templates



HashiCorp
Terraform



Azure Bicep



AWS CloudFormation



Cloud Deployment Manager



ADVANCING ANALYTICS

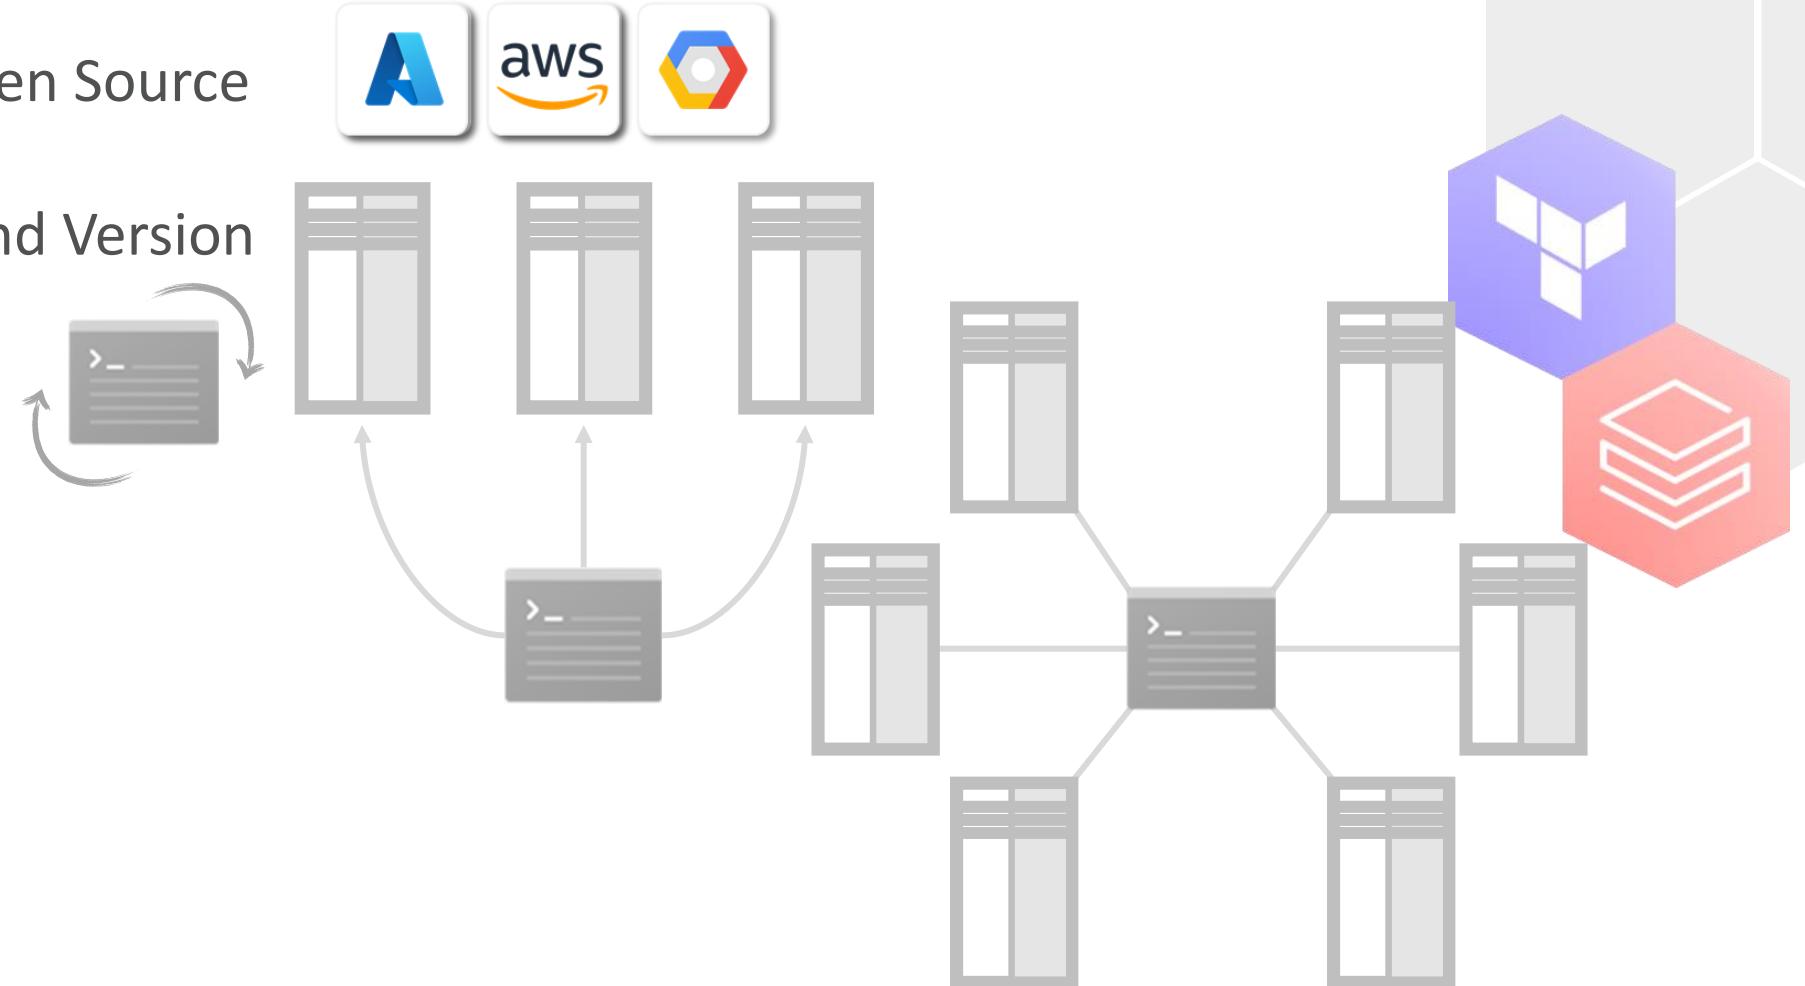


Terraform Overview

Terraform Overview

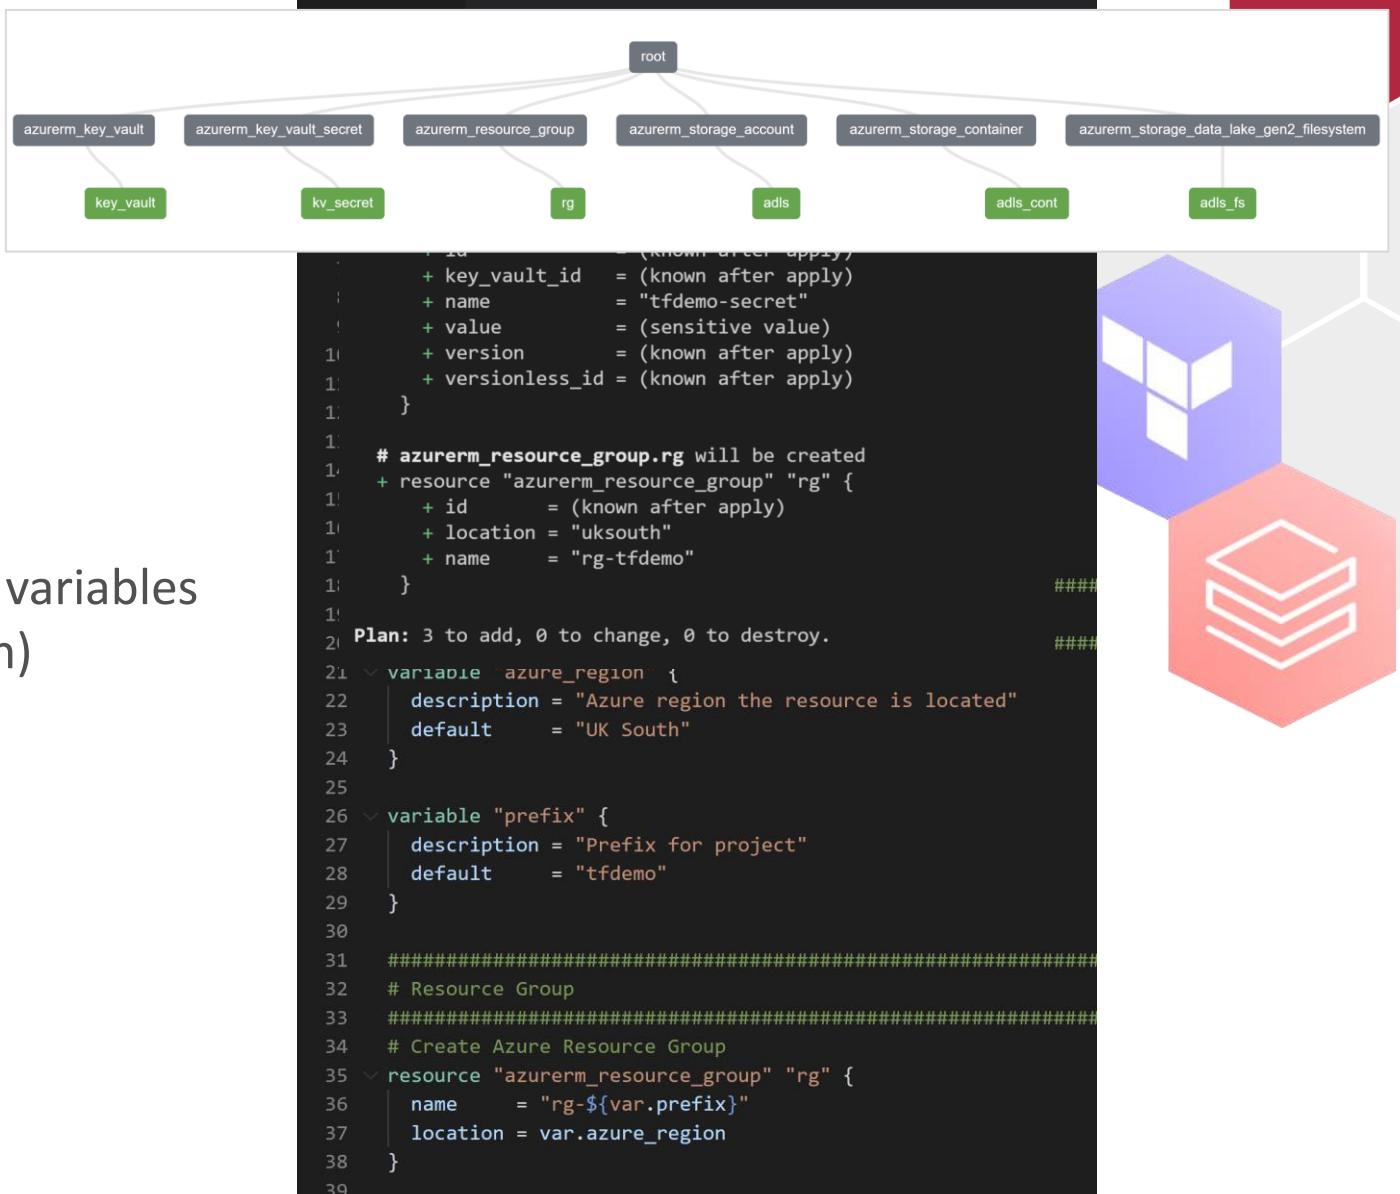


- Infrastructure as Code (IaC)
- Cloud Agnostic and Open Source
- Provisions, Manages and Version
- Consistent
- Reusable



Terraform Features

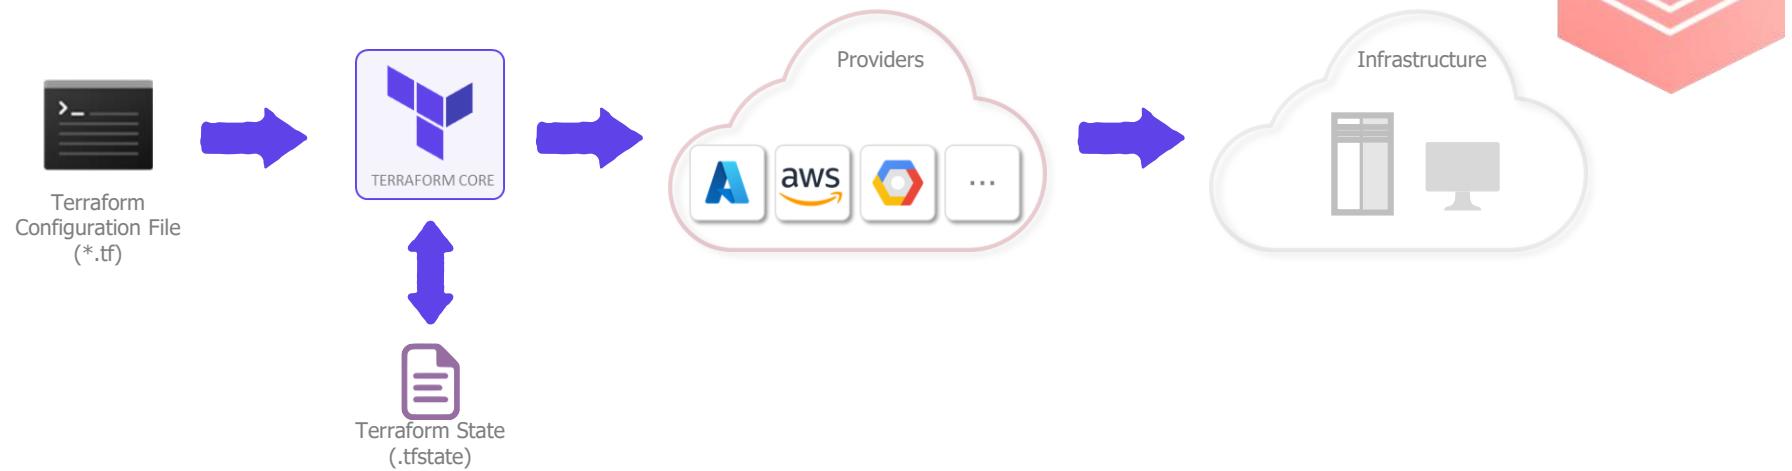
- Resource Graph
- Execution Plan
- Declarative Configure File
- HCL Terraform Language
 - Supports loops, dynamic blocks and local variables
 - Complex data structures (maps, collection)
- Supports JSON



Terraform Architecture



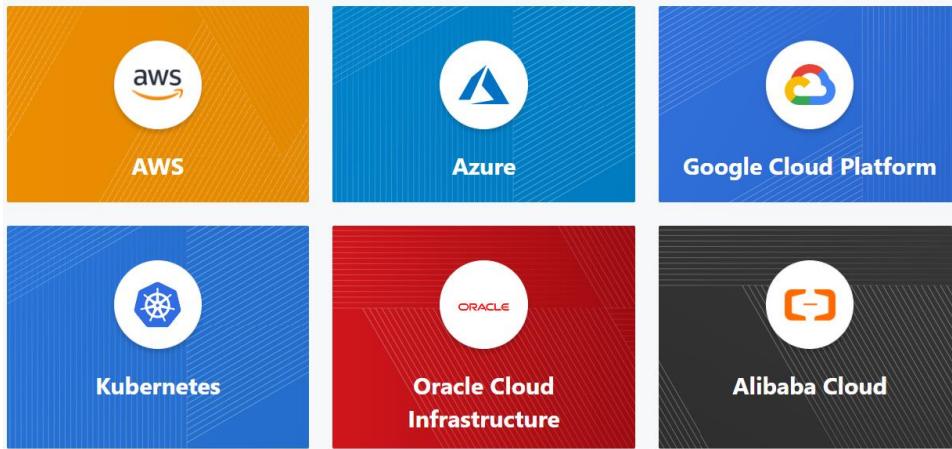
- Configuration File
- Terraform Core Commands
- Providers
- Terraform State
- Infrastructure



Terraform Providers

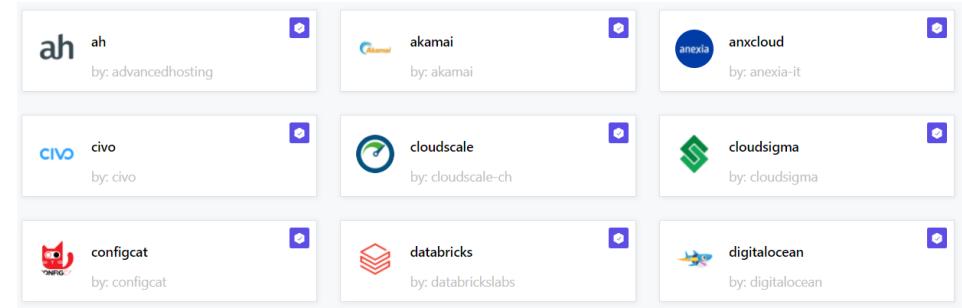
- Providers are required
- Plugins

Official Provider



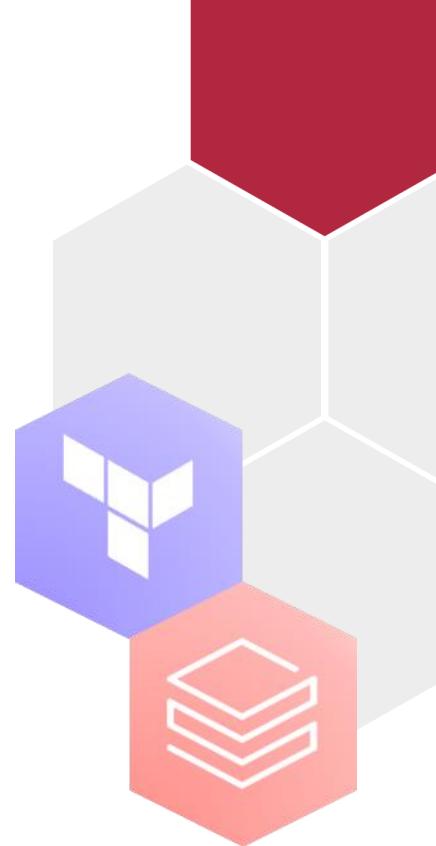
```
terraform {  
  required_providers {  
    azurerm = {  
      source = "hashicorp/azurerm"  
      version = "2.90.0"  
    }  
  }  
  
  provider "azurerm" {  
    # Configuration options  
  }  
}
```

Verified Provider



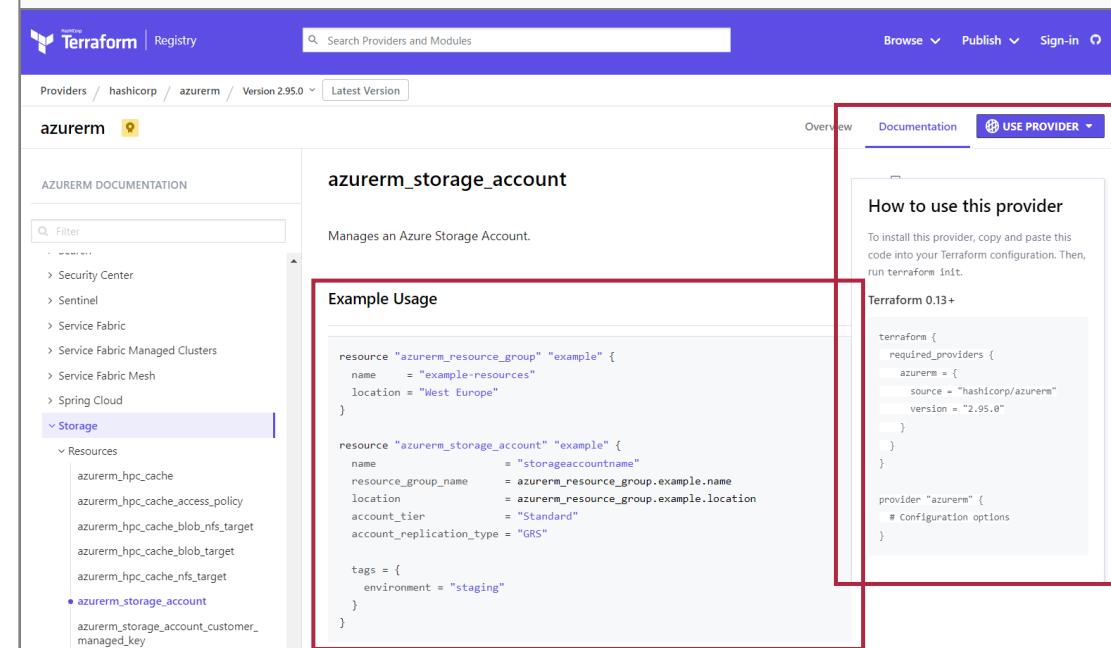
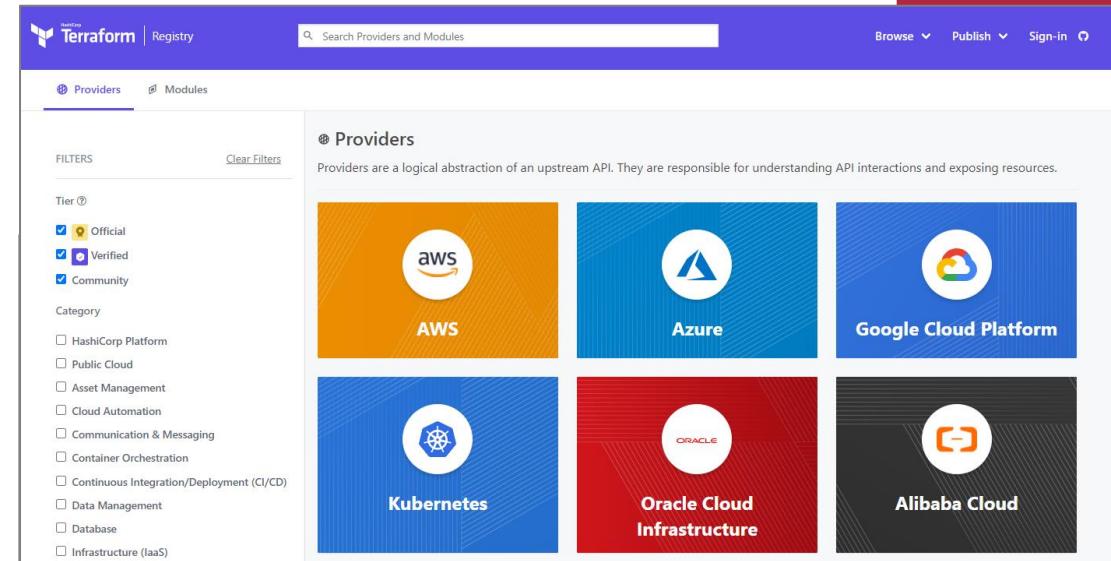
Community Provider

```
terraform {  
  required_providers {  
    azurerm = {  
      source = "hashicorp/azurerm"  
      version = "~> 2.90.0"  
    }  
    databricks = {  
      source = "databrickslabs/databricks"  
      version = "0.4.6"  
    }  
  }  
}
```



Terraform Registry

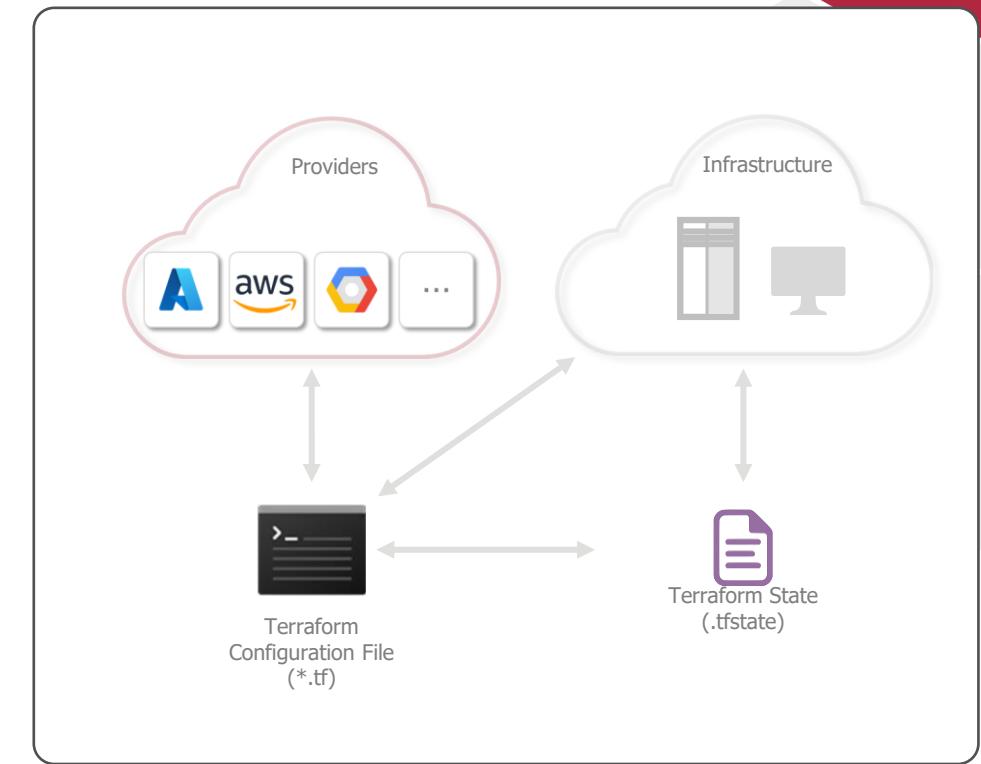
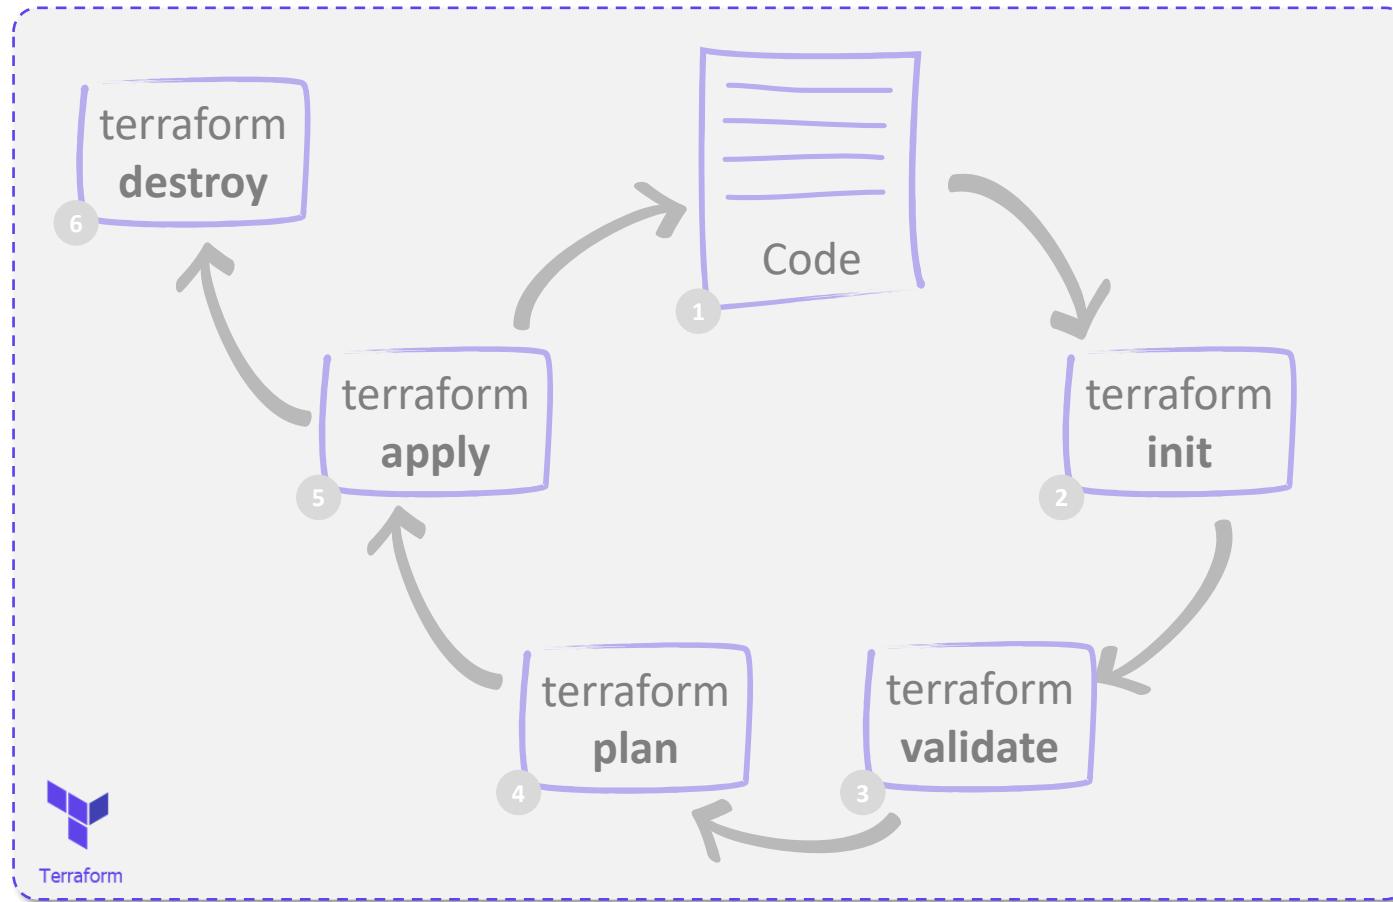
- **Terraform Registry** contains providers and modules. <https://registry.terraform.io/>
- **Example** codes and argument
- **Providers** are plugins that are map to a cloud provider's API
- **Modules** are reusable and self-containing groups of configurations files



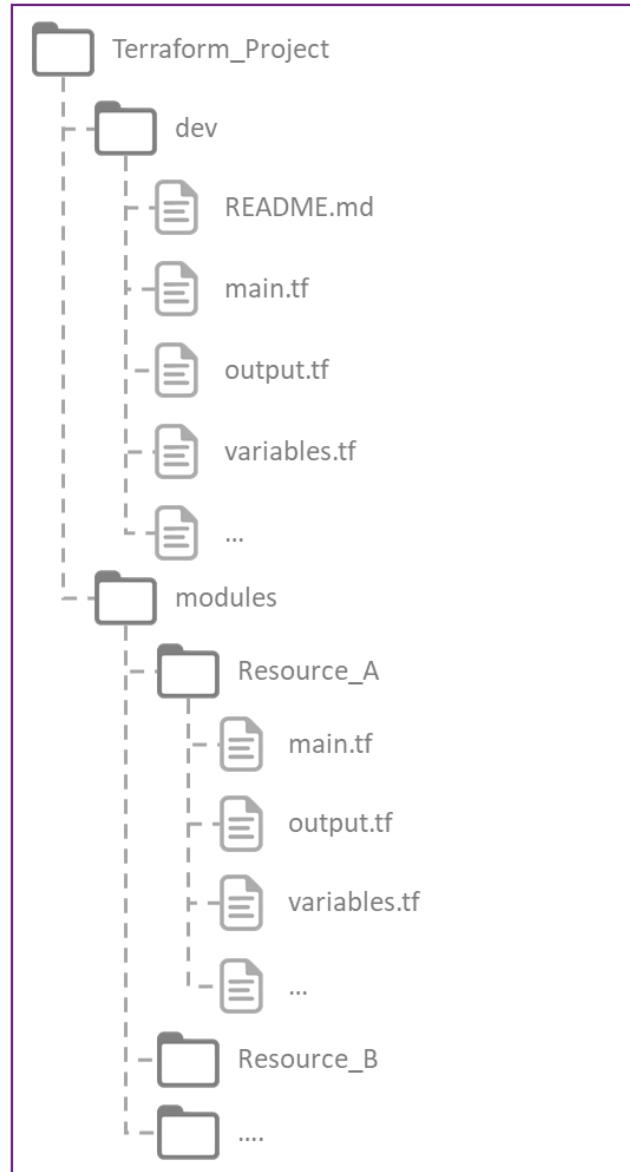


Terraform Concepts

Terraform Lifecycle



Terraform Structure



README.md

- Description of what resources, modules and project

main.tf

- Resource configuration for the infrastructure

output.tf

- Output values of the resources created

variables.tf

- Input variables for resources

*.tfvars

- Variable definitions for the resource to apply

locals.tf

- Define expression and use multiple times within a modules

data.tf

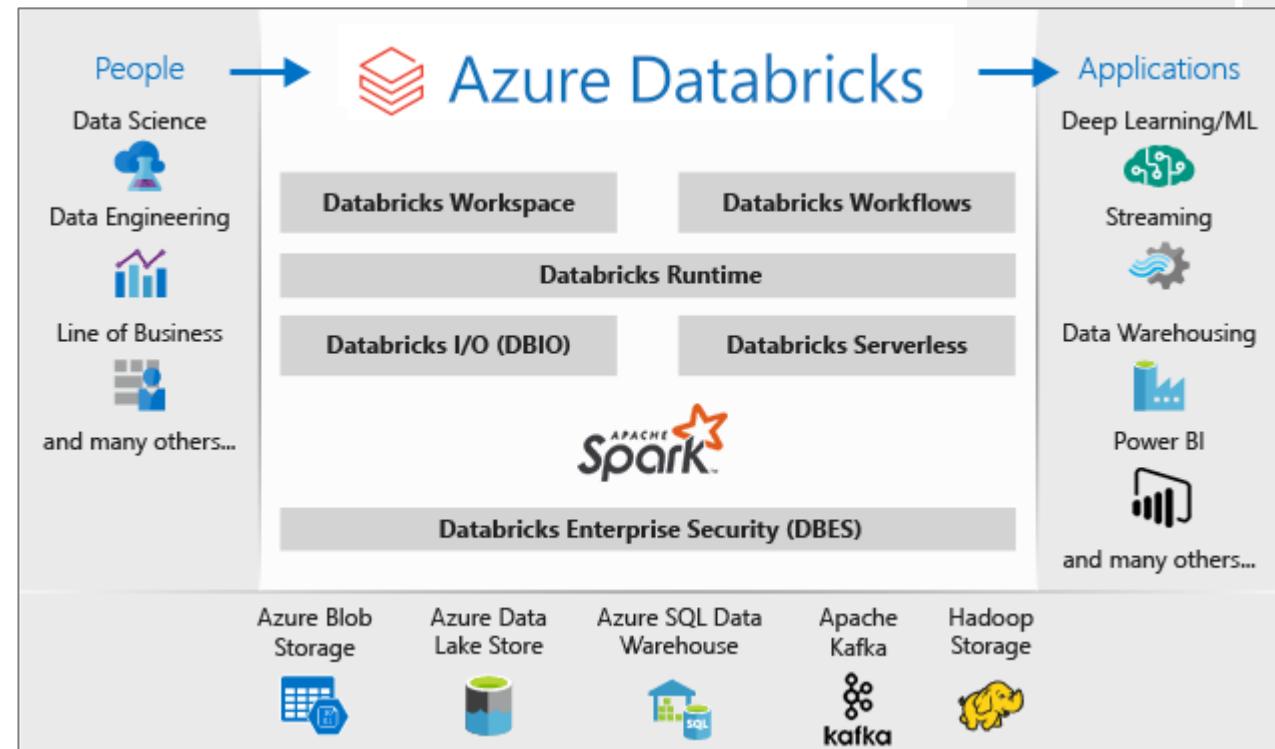
- Return read-only view of pre-existing resource and components





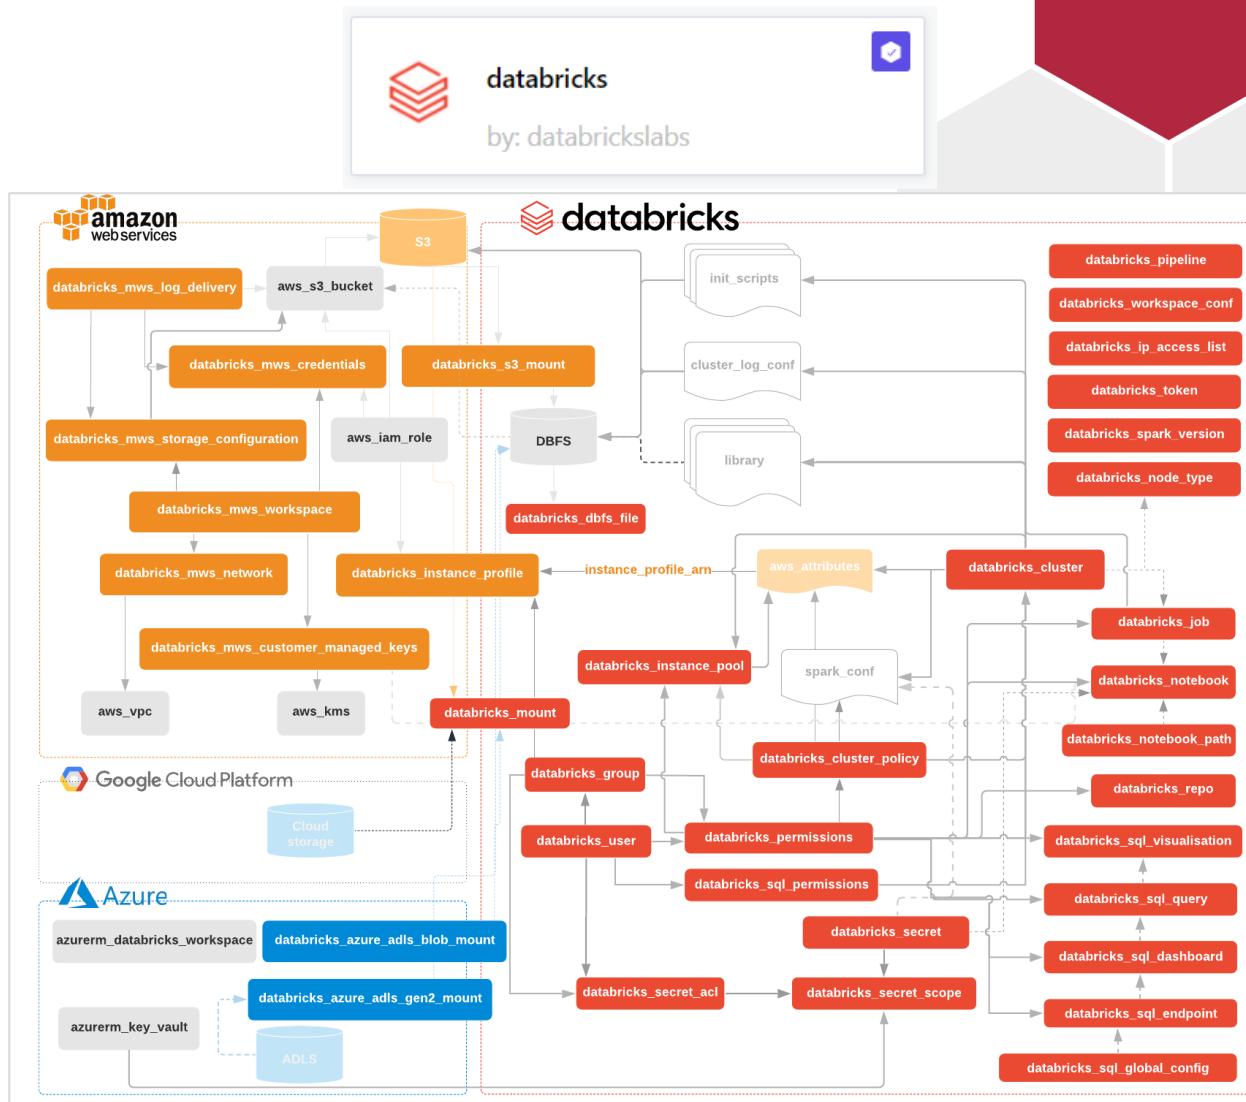
Databricks

- Data analytics platform for cloud services
- Ingest, transform and explore large amounts of data
- Based on Apache-Spark distributed system
 - Multiple Processing
 - Scalable

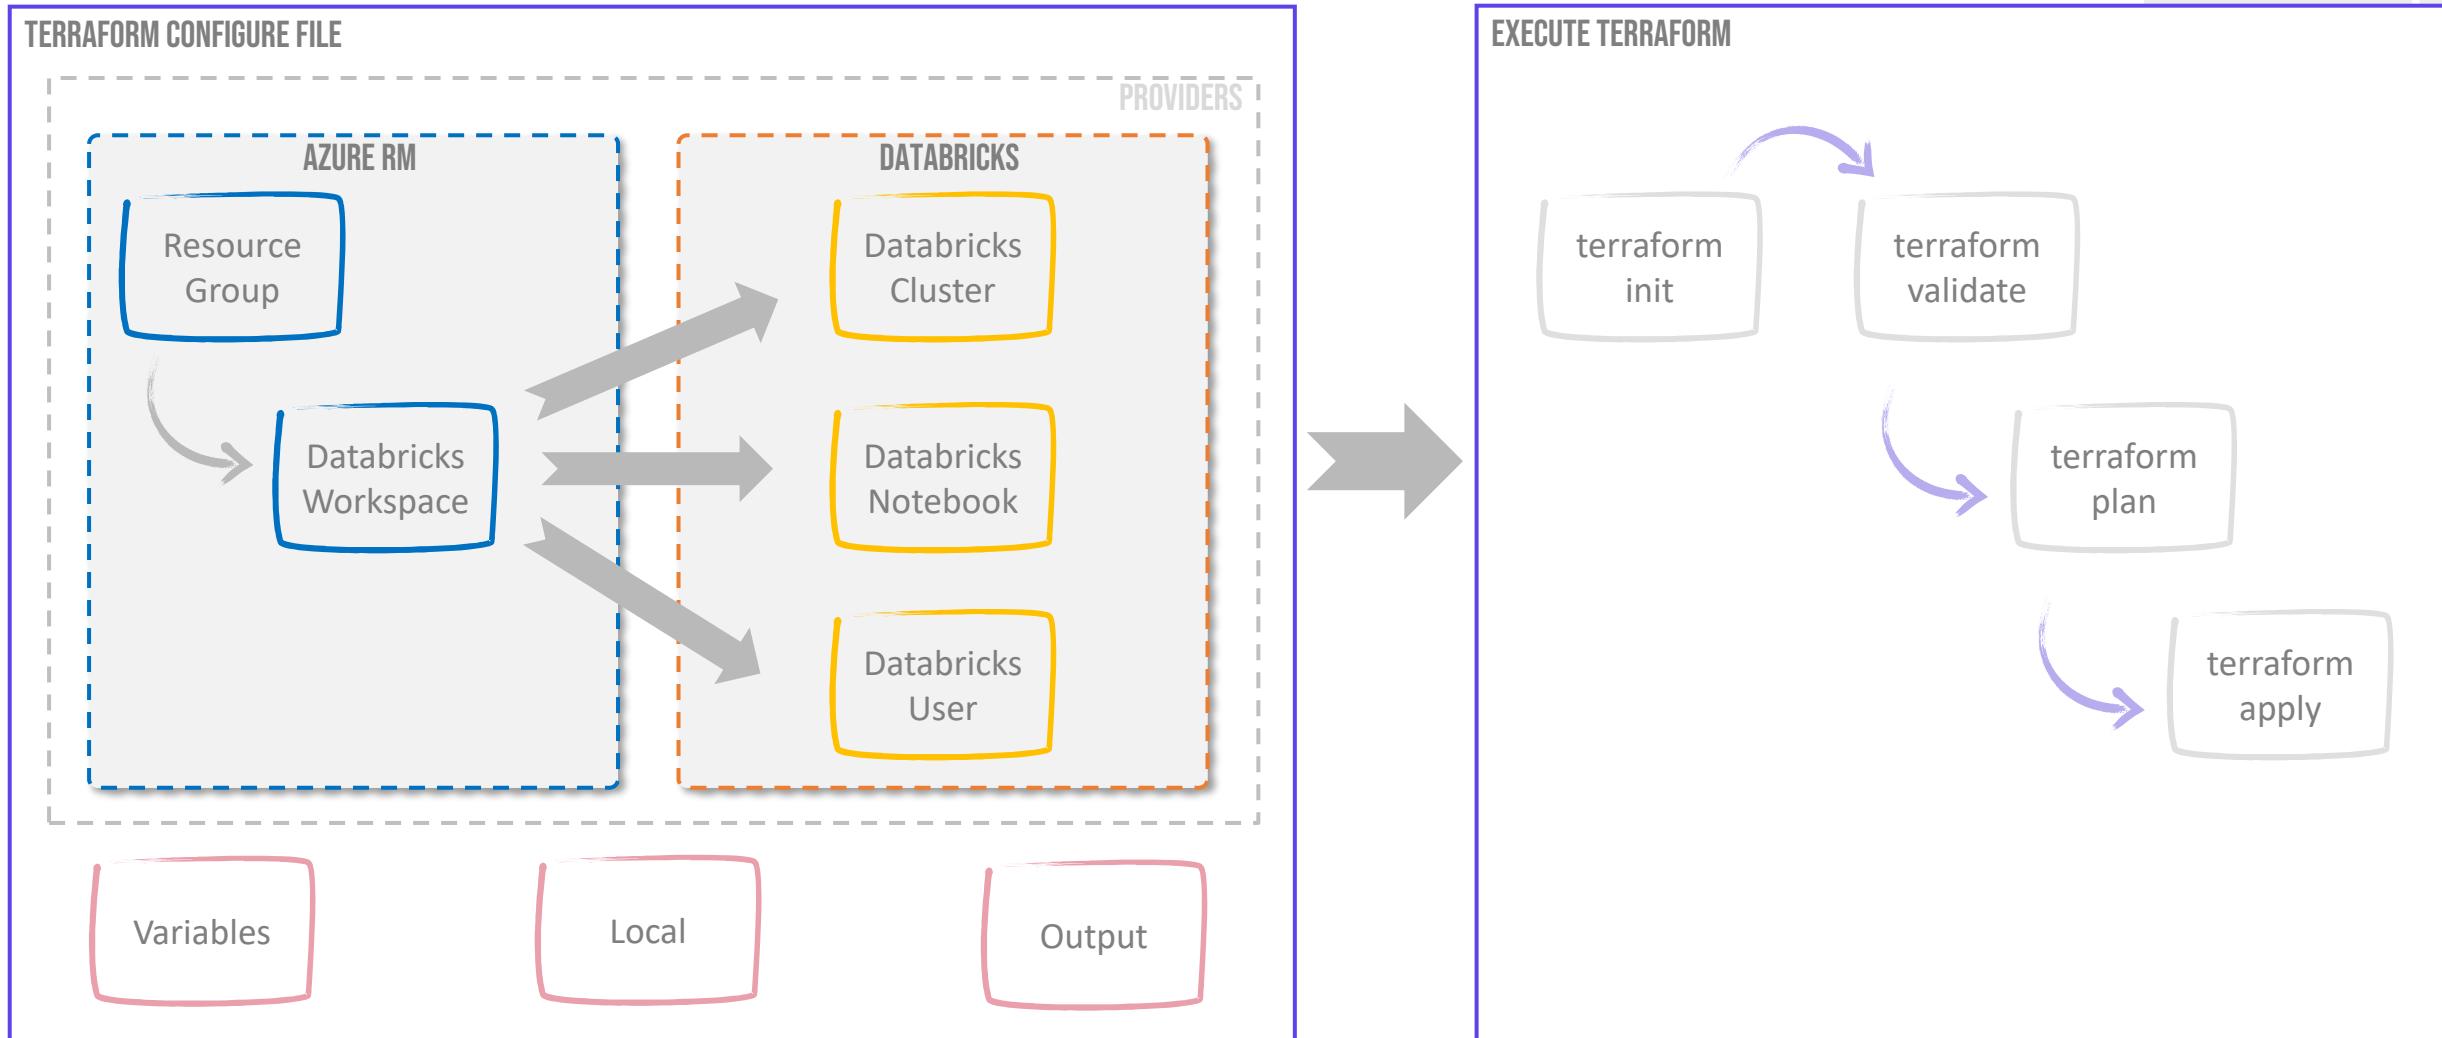


Databricks Labs

- Collection of Terraform providers
- Supports all Databricks REST APIs
- Create, update, and delete components
- Requires both a Cloud Service provider and the Databricks Terraform provider



Deploy Databricks Components using Terraform



Demo

- Using Terraform for Databricks IaC

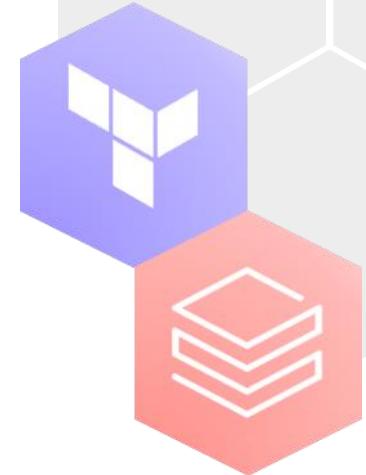




CI/CD Deployment Tools

CI/CD Deployment Tools

- Deploy Terraform using:
 - DevOps CI/CD
 - GitHub Actions
 - Jenkins
 - AWS CodeBuild
 - GitLab
 - CircleCI
 - etc...



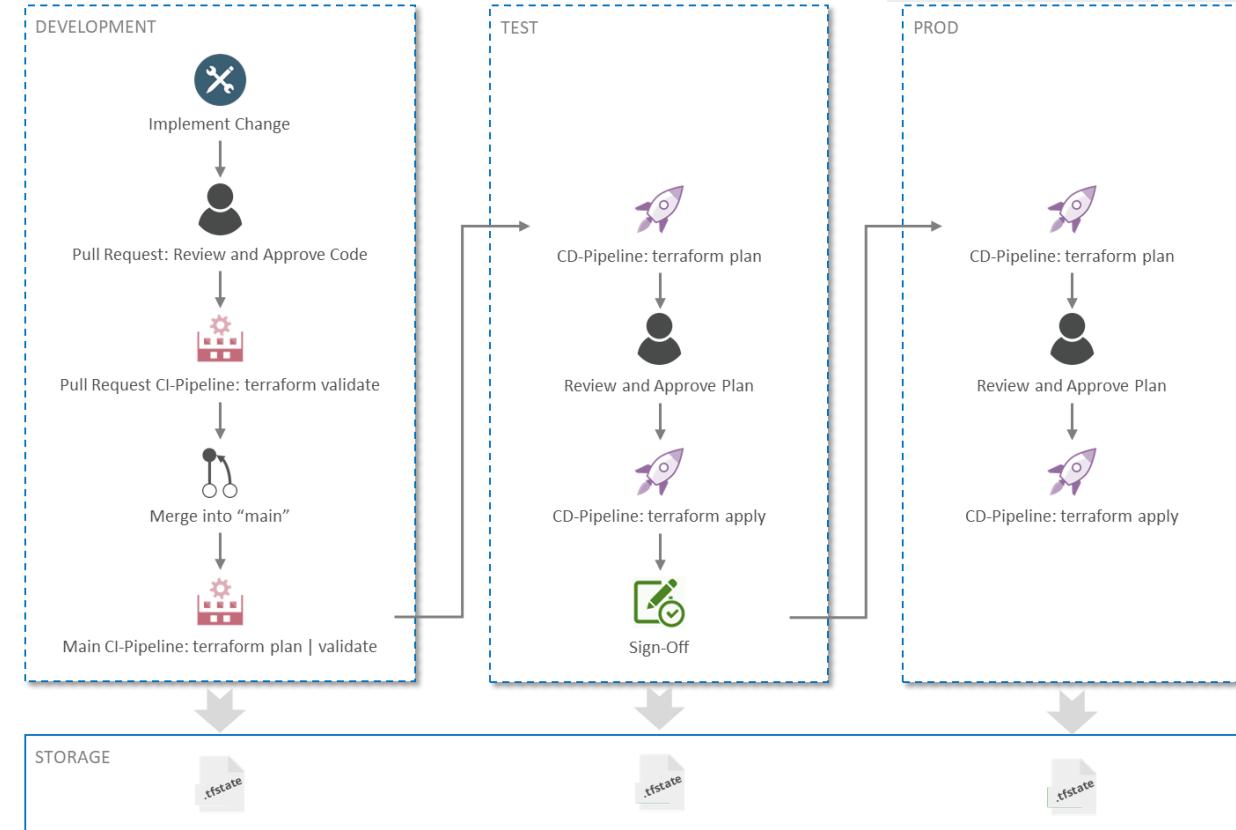
Publish Terraform

Introducing changes into a cloud environment without manual intervention and having explicit permissions

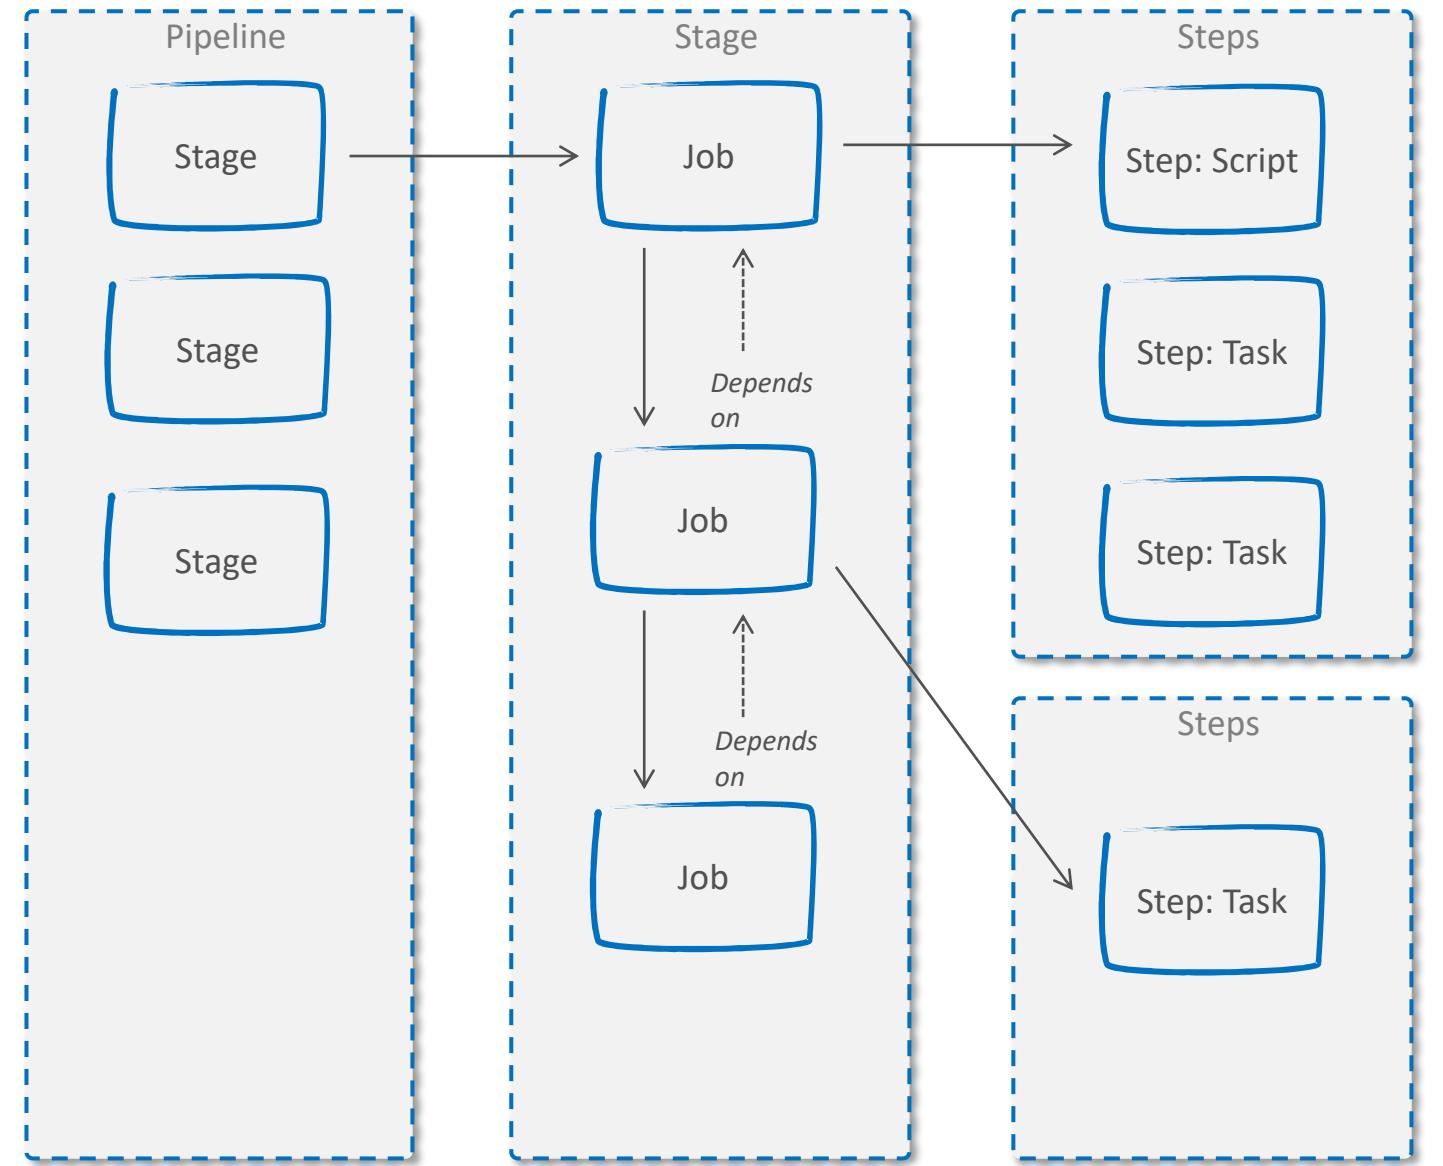
- Develop a deployment pipeline using CI/CD tool for release management
- Adopt GitOps approach and use repositories, branches, and enforce pull requests

Example Delivery Workflow:

- Change is reviewed and merged with a Pull Request
- Pull Request validated using CI-Pipeline
- Branch merged into “main” and validated using CI-Pipeline
- Deploy to Test environment using CD-Pipeline, approve plan and apply
- Unit testing and sign-off environment
- Deploy to Prod environment using CD-Pipeline, approve plan and apply

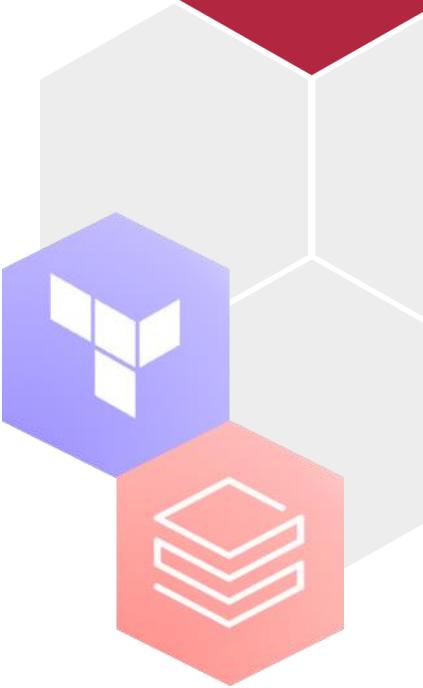


Azure DevOps Pipeline



YML Pipelines consist of:

- Stages
- Jobs
- Steps

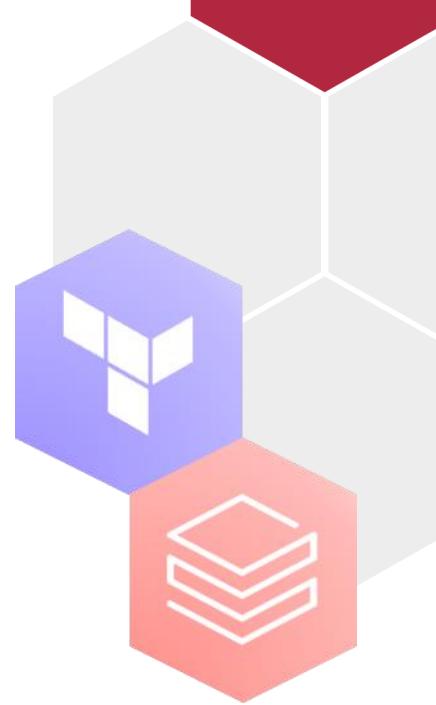
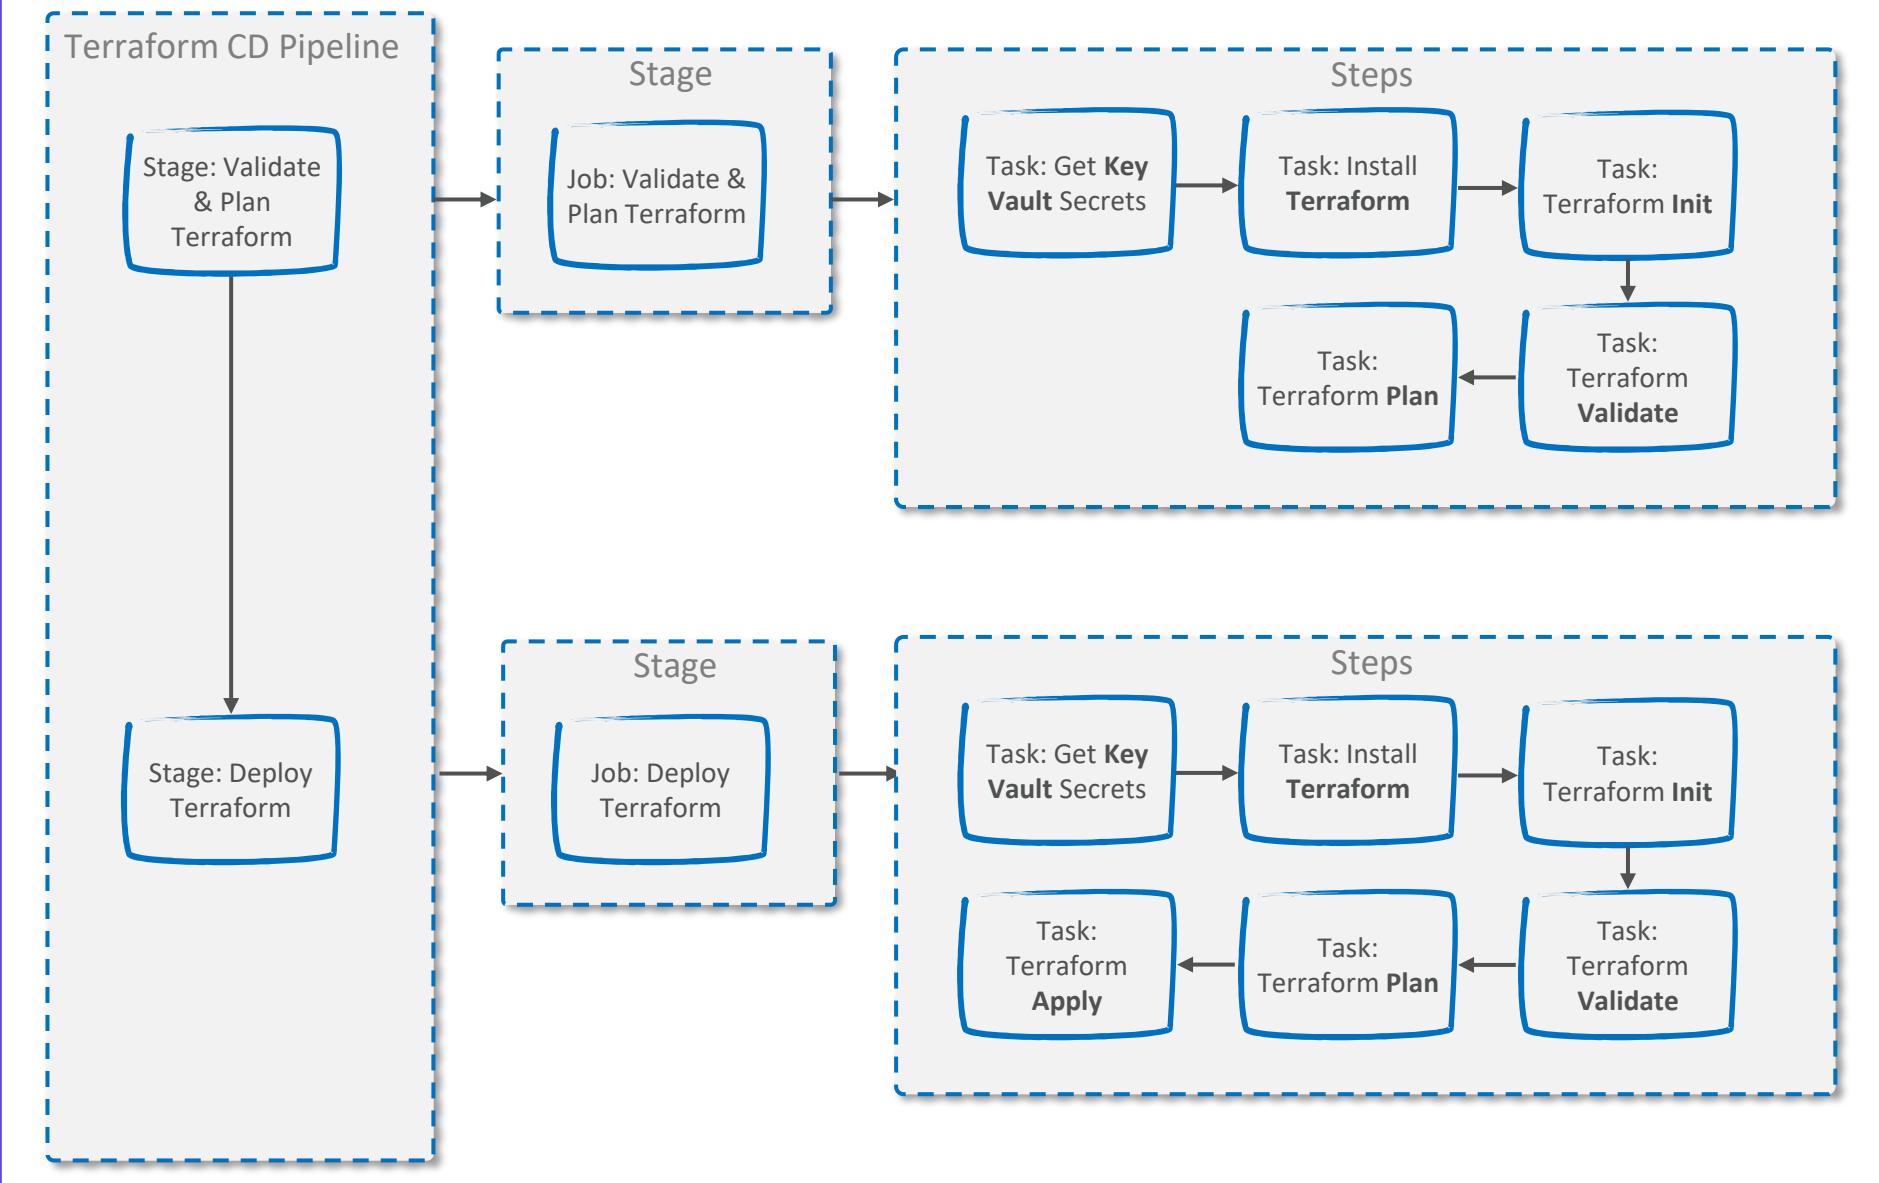


Tasks/Scripts may include:

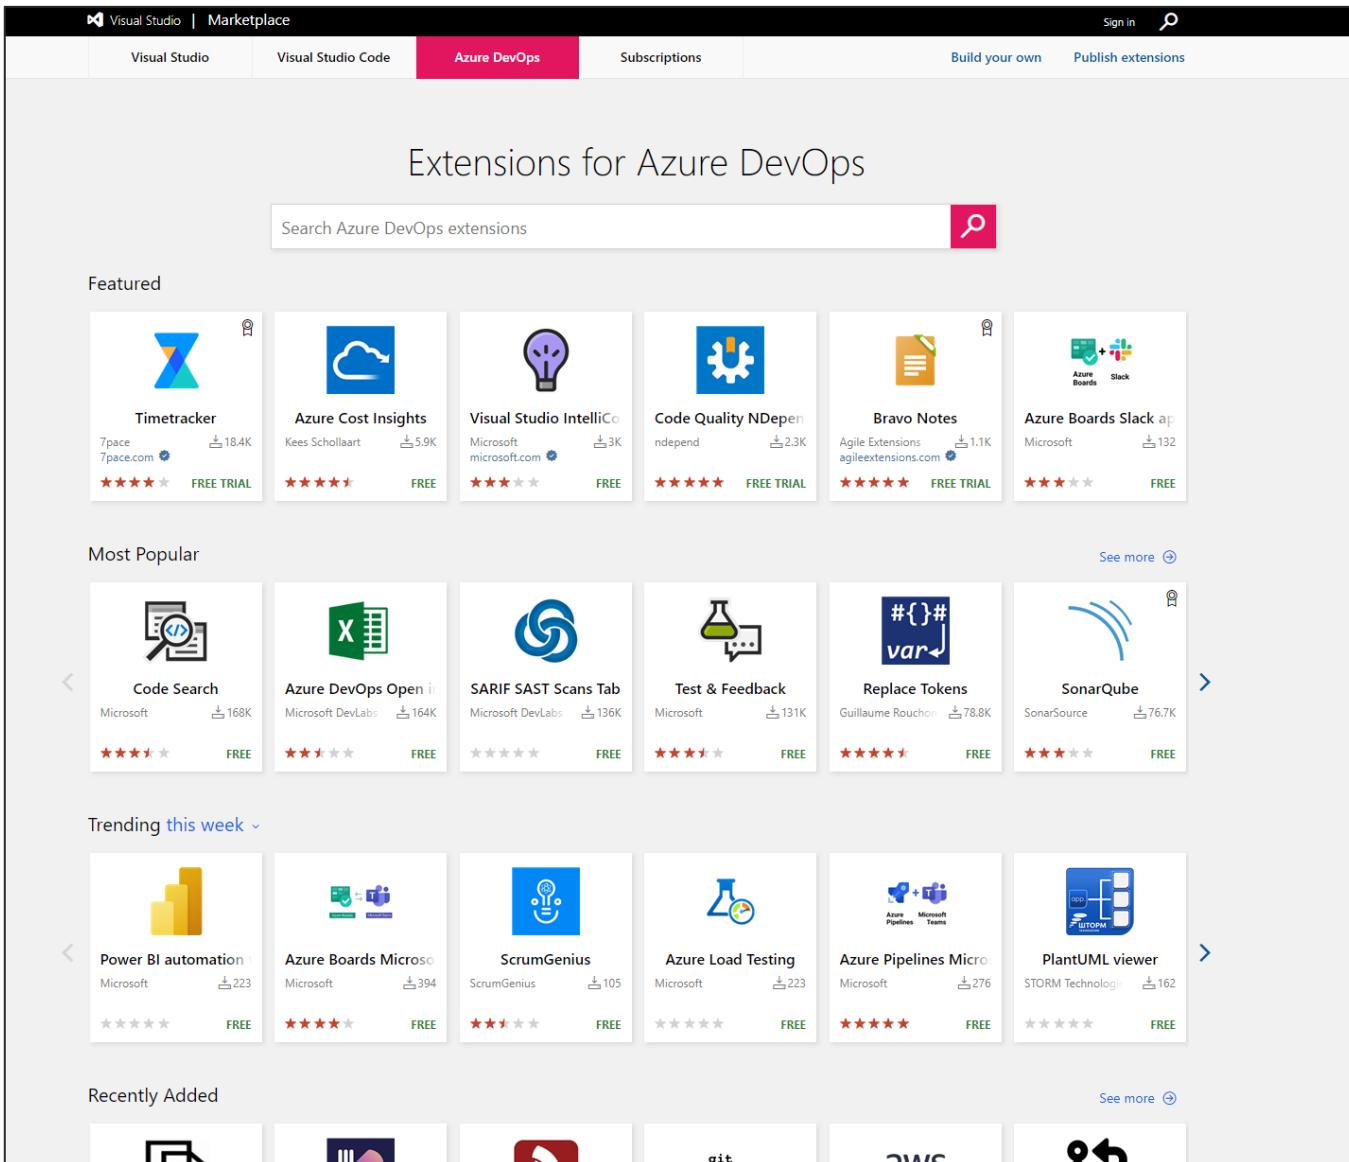
- Install components
- Calls to Rest APIs
- CLI Commands
- Running code (e.g. Python script)
- Connecting to Key Vault
- Publishing Build Artifacts



Azure DevOps Terraform Pipeline



Azure DevOps Extensions



The screenshot shows the Azure DevOps Marketplace interface. At the top, there are tabs for Visual Studio, Visual Studio Code, Azure DevOps (which is highlighted in pink), Subscriptions, and links for Build your own and Publish extensions. A search bar with a magnifying glass icon is positioned above the main content area. The main area is titled "Extensions for Azure DevOps" and features a search bar with the placeholder "Search Azure DevOps extensions".

Featured

- Timetracker** by 7pace (7pace.com) - 18.4K downloads, 4.5 stars, FREE TRIAL
- Azure Cost Insights** by Kees Schollaart - 5.9K downloads, 5 stars, FREE
- Visual Studio IntelliCo** by Microsoft (microsoft.com) - 3K downloads, 4.5 stars, FREE
- Code Quality NDepen** by ndepend - 2.3K downloads, 5 stars, FREE TRIAL
- Bravo Notes** by Agile Extensions (agileextensions.com) - 1.1K downloads, 5 stars, FREE TRIAL
- Azure Boards Slack app** by Microsoft - 132 downloads, 4 stars, FREE

Most Popular

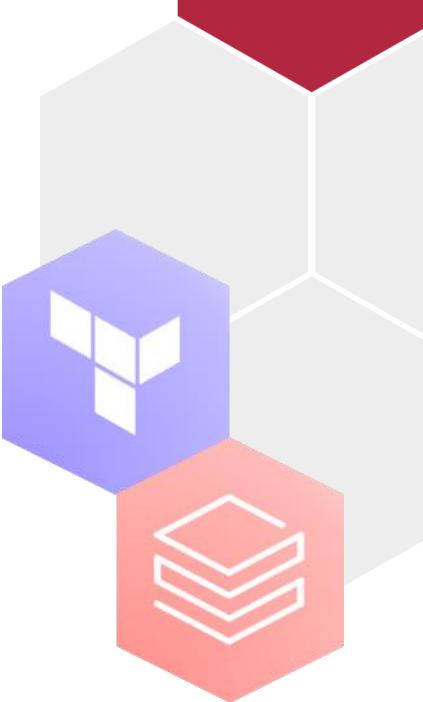
- Code Search** by Microsoft - 168K downloads, 4 stars, FREE
- Azure DevOps Open i** by Microsoft DevLabs - 164K downloads, 4 stars, FREE
- SARIF SAST Scans Tab** by Microsoft DevLabs - 136K downloads, 4 stars, FREE
- Test & Feedback** by Microsoft - 131K downloads, 4 stars, FREE
- Replace Tokens** by Guillaume Rouchon - 78.8K downloads, 4 stars, FREE
- SonarQube** by SonarSource - 76.7K downloads, 4 stars, FREE

Trending this week

- Power BI automation** by Microsoft - 223 downloads, 5 stars, FREE
- Azure Boards Micro** by Microsoft - 394 downloads, 4 stars, FREE
- ScrumGenius** by ScrumGenius - 105 downloads, 4 stars, FREE
- Azure Load Testing** by Microsoft - 223 downloads, 5 stars, FREE
- Azure Pipelines Micro** by Microsoft - 276 downloads, 4 stars, FREE
- PlantUML viewer** by STORM Technologies - 162 downloads, 5 stars, FREE

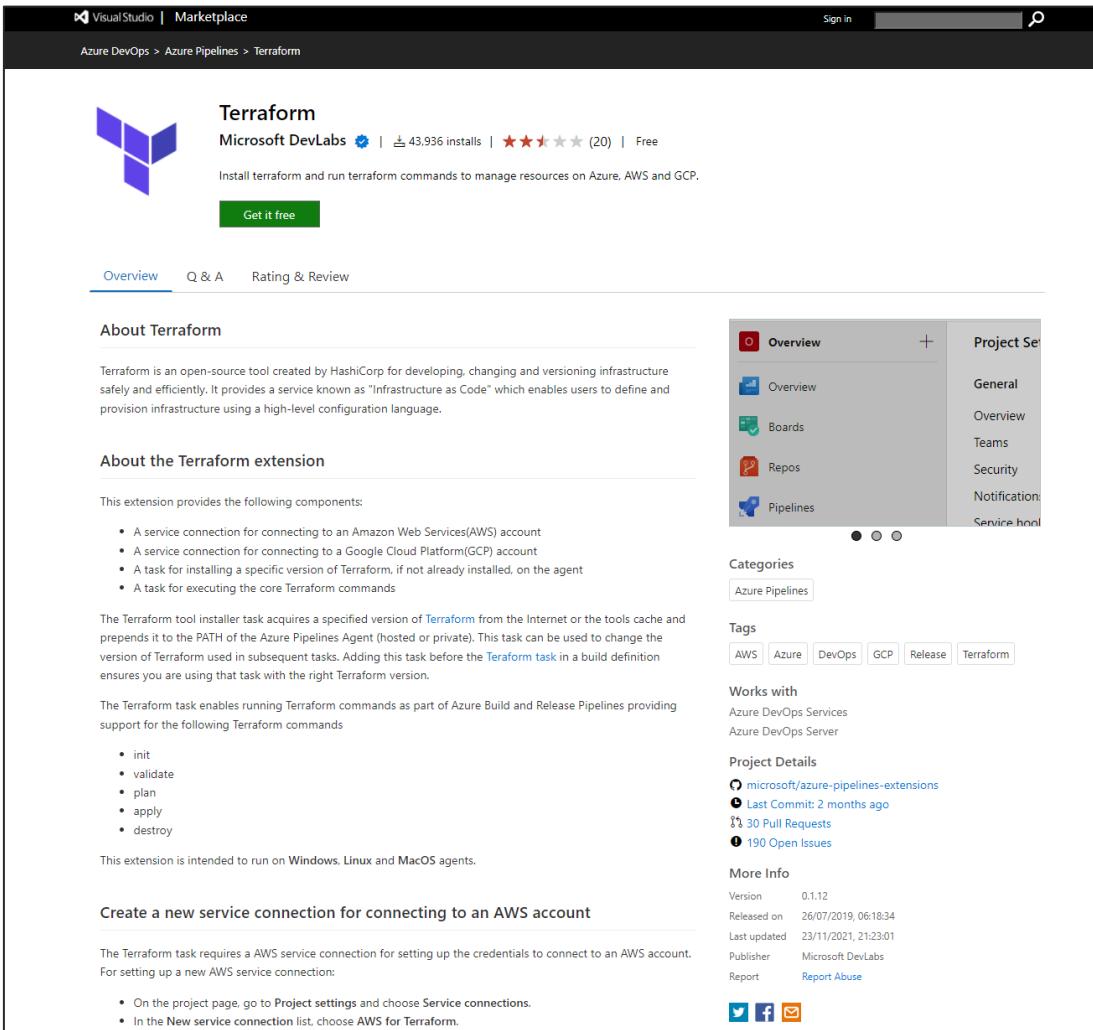
Recently Added

- Recent extensions shown include a dashboard icon, a chart icon, a file icon, git, AWS, and a key icon.



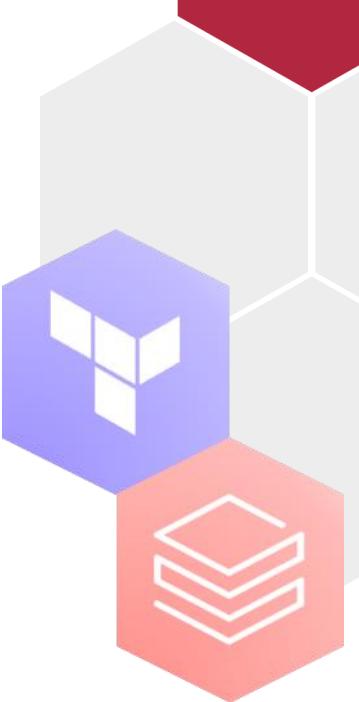
Azure DevOps Terraform Extension

Recommended:



The screenshot shows the Azure DevOps Terraform Extension page in the Visual Studio Marketplace. The page includes the following sections:

- Overview:** Shows the extension has 43,936 installs and a 4-star rating from 20 reviews. A "Get it free" button is present.
- About Terraform:** Describes Terraform as an open-source tool for infrastructure as code.
- About the Terraform extension:** Lists components: AWS and GCP service connections, Terraform installation task, and Terraform command execution task.
- Tooling:** Describes the Terraform tool installer task and its use in Azure Pipelines.
- Usage:** Lists Terraform commands supported by the extension: init, validate, plan, apply, and destroy.
- System Requirements:** States the extension is intended to run on Windows, Linux, and MacOS agents.
- Create a new service connection for connecting to an AWS account:** Instructions for setting up an AWS service connection.
- Project Details:** Includes links to the repository on GitHub, last commit date (2 months ago), 30 pull requests, and 190 open issues.
- More Info:** Lists version (0.1.12), release date (26/07/2019), update date (23/11/2021), publisher (Microsoft DevLabs), and a "Report Abuse" link.



ADVANCING ANALYTICS

Getting started

pipeline.yml

No trigger

Parameters entered on pipeline execution

Variables: static, dynamic and key Vault

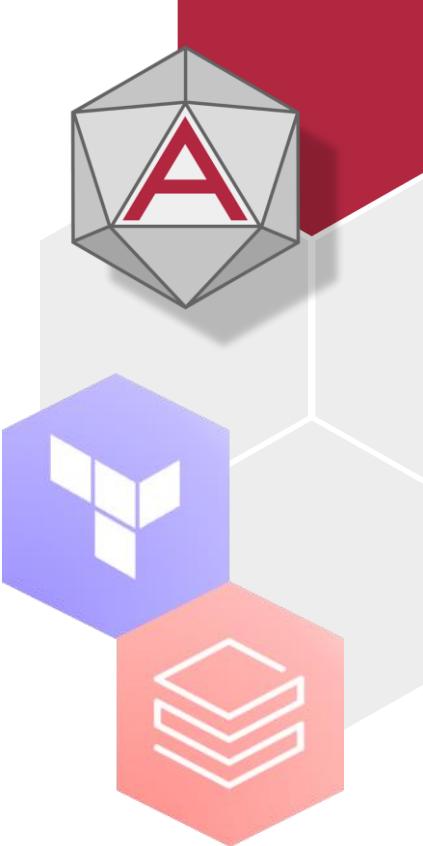
```
trigger: none

parameters:
  - name: par_environment
    displayName: Enter the Environment Name
    default: dev
    type: string

variables:
  env: $(env)
  terraform_directory: tf05_publish-terraform-ci-cd
  terraform_version: latest
  service_connection:
    key_vault_name:
    tf_state_backend_resource_group_name:
    tf_state_backend_resource_group_location: 'UK South'
    tf_state_backend_storage_account_name:
    tf_state_backend_container_name: terraform-state-$(env)
    tf_state_backend_key_name: $(env).terraform.tfstate

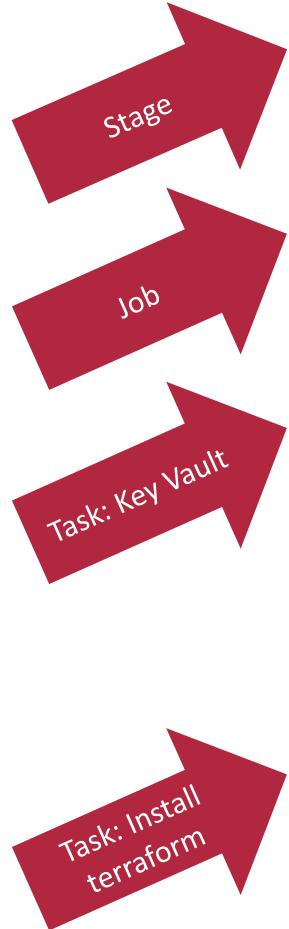
pool:
  vmImage: 'ubuntu-latest'
```

Microsoft hosted image for pool





Stage



```
stages :  
- stage: Validate_Plan_Terraform  
  jobs:  
    - job: Validate_Plan_Terraform  
      displayName: "Validate & Plan Terraform > install, init, validate and plan"  
      continueOnError: false  
      steps:  
        - checkout: self  
        - task: AzureKeyVault@1  
          displayName: Retrieve key vault secrets  
          inputs:  
            azureSubscription: $(service_connection)  
            keyVaultName: $(key_vault_name)  
            secretsFilter: 'ARM-CLIENT-ID, ARM-CLIENT-SECRET, ARM-TENANT-ID, ARM-SUBSCRIPTION-ID'  
            runAsPreJob: false  
  
    - task: TerraformInstaller@0  
      displayName: Install Terraform  
      inputs:  
        terraformVersion: $(terraform_version)
```



ADVANCING ANALYTICS



Init

Terraform task

```
- task: TerraformCLI@0
  displayName: Terraform Init
  inputs:
    command: "init"
    workingDirectory: $(System.DefaultWorkingDirectory)/$(terraform_directory)
    backendType: "azurerm"
    allowTelemetryCollection: true
    backendServiceArm: $(service_connection)
    runAzLogin: true
    ensureBackend: true
    backendAzureRmResourceGroupName: $(tf_state_backend_resource_group_name)
    backendAzureRmResourceGroupLocation: $(tf_state_backend_resource_group_location)
    backendAzureRmStorageAccountName: $(tf_state_backend_storage_account_name)
    backendAzureRmContainerName: $(tf_state_backend_container_name)
    backendAzureRmKey: '$(tf_state_backend_key_name)'
  env:
    ARM_CLIENT_ID: $(ARM-CLIENT-ID)
    ARM_CLIENT_SECRET: $(ARM-CLIENT-SECRET)
    ARM_SUBSCRIPTION_ID: $(ARM-SUBSCRIPTION-ID)
    ARM_TENANT_ID: $(ARM-TENANT-ID)
```

State Configuration



ADVANCING ANALYTICS

Apply



Changed the command

```
- task: TerraformCLI@0
  displayName: Terraform Apply
  inputs:
    command: "apply"
    workingDirectory: $(System.DefaultWorkingDirectory)/$(terraform_directory)
    backendType: "azurerm"
    allowTelemetryCollection: true
    backendServiceArm: $(service_connection)
    runAzLogin: true
    ensureBackend: true
    backendAzureRmResourceGroupName: $(tf_state_backend_resource_group_name)
    backendAzureRmResourceGroupLocation: $(tf_state_backend_resource_group_location)
    backendAzureRmStorageAccountName: $(tf_state_backend_storage_account_name)
    backendAzureRmContainerName: $(tf_state_backend_container_name)
    backendAzureRmKey: '$(tf_state_backend_key_name)'
  env:
    ARM_CLIENT_ID: $(ARM-CLIENT-ID)
    ARM_CLIENT_SECRET: $(ARM-CLIENT-SECRET)
    ARM_SUBSCRIPTION_ID: $(ARM-SUBSCRIPTION-ID)
    ARM_TENANT_ID: $(ARM-TENANT-ID)
```

Demo

- Running a basic pipeline
- Implementing gating





Recap

Recap



Deploy Quickly
and Effectively



Easy to Read,
Manage and
Maintain



Faster
Development



Repeatable



Consistent



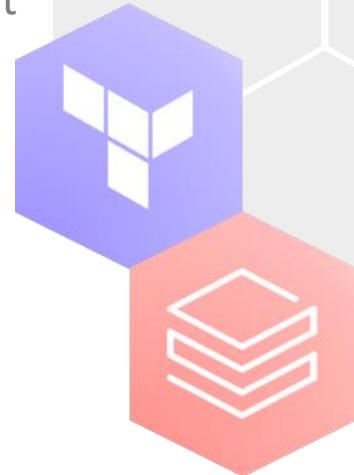
Various CI/CD
Tools



Governance and
Controls



GitOps



TRY IT OUT



databricks



HashiCorp
Terraform



Azure DevOps



Check out our blogs on Databricks, Terraform and DevOps



ADVANCING ANALYTICS

<https://sqlb.it/?7101>



 [@annawykes](https://twitter.com/annawykes)
 <https://www.linkedin.com/in/anna-maria-wykes-31939454>
 [AnnaWykes](https://github.com/AnnaWykes)



 [@falekmiah](https://twitter.com/falekmiah)
 [falekmiah.com](https://github.com/falekmiah.com)
 [FalekMiah01](https://github.com/FalekMiah01)

Code base can be found on GitHub:
<https://github.com/FalekMiah01/sqlbits2022>

Deploy Databricks Components using Terraform

